



Discrete Techniques

The Computer Graphics Pipeline

- Fragment processing
 - Operations to determine final pixel colour
 - Buffer operations
 - Colour buffer, depth buffer, stencil buffer, etc.
 - Texture mapping
 - Blending
 - Per-pixel lighting





What are Fragments?

- After surface rasterized, no longer notion of a polygon
 - Surface essentially been “broken up” or “discretised” into small pieces
 - Each of which is at most the size of one pixel, called fragments
 - Picture elements, or pixels, are the final screen space colour of the rendered image
 - Conceptually fragments can be smaller than one pixel
 - More than one fragment can contribute to the colour of a pixel


Mapping Techniques

- Texture mapping
 - There are limitation to geometric modeling
 - “Paint” image onto polygons
 - Analogy – sticking wallpaper onto a wall
 - Increases the visual realism



Mapping Techniques

- Texture mapping



Model

Model + Shading



Model + Shading
+ Textures


At what point
do things start
looking real?



For more info on the computer artwork of Jeremy Birn
see <http://www.3drender.com/jbirn/productions.html>




Mapping Techniques

- Texture mapping




Mapping Techniques

- Texture mapping
 - Different texture maps can be used for same object



Mapping Techniques


- Texture mapping




Mapping Techniques

- Texture mapping


➤ Texture coordinates are referred to as (s, t)







Texture



Polygon



Mapping Techniques



Mapping Techniques

➤ Texture coordinates

```
struct VertexTex {  
    GLfloat position[3];  
    GLfloat texCoord[2];  
};  
  
std::vector<GLfloat> gVertices = {  
    -0.2f, -0.5f, 0.0f, // vertex 0: position  
    0.0f, 0.0f, // vertex 0: texture coordinate  
    0.8f, -0.5f, 0.0f, // vertex 1: position  
    1.0f, 0.0f, // vertex 1: texture coordinate  
    -0.2f, 0.5f, 0.0f, // vertex 2: position  
    0.0f, 1.0f, // vertex 2: texture coordinate  
    0.8f, 0.5f, 0.0f, // vertex 3: position  
    1.0f, 1.0f, // vertex 3: texture coordinate  
};
```

Mapping Techniques

➤ Texture coordinates

```
glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE,  
    sizeof(VertexTex),  
    reinterpret_cast<void*>(offsetof(VertexTex, position)) );  
glVertexAttribPointer(1, 2, GL_FLOAT, GL_FALSE,  
    sizeof(VertexTex),  
    reinterpret_cast<void*>(offsetof(VertexTex, texCoord)) );  
  
 glEnableVertexAttribArray(0);  
 glEnableVertexAttribArray(1);
```

Mapping Techniques


➤ Texture coordinates

- In vertex shader

```
layout(location = 0) in vec3 aPosition;  
layout(location = 1) in vec2 aTexCoord;  
  
out vec2 vTexCoord;  
  
void main()  
{  
    ...  
    // interpolate texture coordinate  
    vTexCoord = aTexCoord;  
}
```

Mapping Techniques

- Texture mapping
 - Texture coordinates and wrapping



Mapping Techniques

- Texture mapping

- Wrapping

- Texture coordinates outside the 0.0 to 1.0 range
- GL_REPEAT : Default, repeats the image
- GL_MIRRORED_REPEAT : Mirrors each repeat
- GL_CLAMP_TO_EDGE : Use edge colour
- GL_CLAMP_TO_BORDER : Use a user defined border colour




GL_REPEAT



GL_MIRRORED_REPEAT



GL_CLAMP_TO_EDGE



GL_CLAMP_TO_BORDER

Mapping Techniques


➤ Wrapping

```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S,  
    GL_CLAMP_TO_EDGE);
```


```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T,  
    GL_REPEAT);
```

- If clamping to border colour


```
float borderColor[] = { 1.0f, 0.0f, 0.0f, 1.0f };  
glTexParameterfv(GL_TEXTURE_2D,  
    GL_TEXTURE_BORDER_COLOR, borderColor);
```



texture




GL_REPEAT
wrapping



GL_CLAMP_TO_EDGE
wrapping




Mapping Techniques

- Texture mapping
 - Aliasing in textures



Mapping Techniques


- Texture mapping
 - Aliasing is the under-sampling of a signal
 - In texture mapping, aliasing artifacts emerge as a result of discrete sampling
 - Looks worse when moving



Mapping Techniques


- Texture mapping
 - Size of pixel we are trying to colour may be larger or smaller than one texture element (**texel**)

Magnification



(a)

Minification





(b)

Mapping Techniques

- Texture mapping

- Filtering

- GL_NEAREST : Default filtering method
 - Nearest neighbour filtering
 - » Select colour closest to texture coordinate
- GL_LINEAR
 - (bi)linear filtering
 - » Interpolate value from neighbouring texels
 - » The closer to a colour, the greater its contribution





Mapping Techniques

- Texture mapping

- Mipmaps

- For minification problem
 - Create multiple resolutions of an image
 - E.g., for a 256 x 256 image
 - » Create 128x128, 64x64, 32x32, 16x16, 8x8, 4x4, 2x2, 1x1
 - Sample used depends on polygon's size on screen





Mapping Techniques

➤ Mipmap filtering

- Artifacts like sharp edges may be visible when switching between mipmap layers
- Filter between mipmap levels
 - `GL_NEAREST_MIPMAP_NEAREST` : Nearest mipmap to match the pixel size and nearest neighbor interpolation for texture sampling
 - `GL_LINEAR_MIPMAP_NEAREST` : Nearest mipmap level and samples that level using linear interpolation
 - `GL_NEAREST_MIPMAP_LINEAR` : Linearly interpolates between the two mipmaps that most closely match the size of a pixel and samples the interpolated level via nearest neighbor interpolation
 - `GL_LINEAR_MIPMAP_LINEAR` : Linearly interpolates between the two closest mipmaps and samples the interpolated level via linear interpolation

Mapping Techniques

- Texture mapping
 - Mipmaps can be used in conjunction with texture filtering methods



Mapping Techniques

- Generating a texture

 - Using an image library

```
#define STB_IMAGE_IMPLEMENTATION  
#include "stb_image.h"
```

 - Load image data

```
stbi_set_flip_vertically_on_load(true);  
  
int imageWidth, imageHeight, imageChannels;  
unsigned char* imageData =  
    stbi_load("./images/check.bmp",  
    &imageWidth, &imageHeight, &imageChannels, 0);
```

Mapping Techniques

- Generating a texture

- Generate OpenGL texture object

```
GLuint gTextureID;
```

```
...
```

```
glGenTextures(1, &gTextureID);
```

```
glBindTexture(GL_TEXTURE_2D, gTextureID);
```

- Create texture and mipmaps

```
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGB, imageWidth,  
            imageHeight, 0, GL_RGB, GL_UNSIGNED_BYTE,  
            imageData);
```


```
glGenerateMipmap(GL_TEXTURE_2D);
```

Mapping Techniques


- Generating a texture
 - Set texture parameters

```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAX_FILTER,  
                GL_NEAREST);
```

```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER,  
                GL_NEAREST_MIPMAP_NEAREST);
```



Texture
Magnification



Texture
Minification

Mapping Techniques

- Using the texture

➤ In fragment shader

```
in vec2 vTexCoord;
```

```
uniform sampler2D uTextureSampler;
```

```
out vec3 fColor;
```

```
void main()
```

```
{
```

```
    fColor = texture(uTextureSampler, vTexCoord).rgb;
```

```
}
```

Mapping Techniques

- Using the texture

 - In the render_scene() function



```
gShader.use();
```

```
// set texture
gShader.setUniform("uTextureSampler", 0);
glActiveTexture(GL_TEXTURE0);
glBindTexture(GL_TEXTURE_2D, gTextureID);

glBindVertexArray(gVAO);
glDrawArrays(GL_TRIANGLE_STRIP, 0, 4);
```

Mapping Techniques

- Multi-texturing
 - Paint multiple textures onto the same surface




Mapping Techniques

- Light mapping
 - Illumination maps/light maps were introduced in the game *Quake* to represent lighting effects



Mapping Techniques

- Light mapping
 - Allows lighting to be pre-calculated and stored as 2D texture map



Quake 2's packed
light map texture




Mapping Techniques

- Light mapping



X



≡

Quake 2 light mapping




Buffers

- Per-pixel data is stored in buffers
 - Colour Buffer (front and back)
 - Depth Buffer
 - Stencil Buffer
 - Others (overlay planes, auxiliary buffers, colour indices)


Buffers

- The colour buffer
 - Where the RGBA pixel values are stored
 - Generally need two separate colour buffers
 - One for displaying, the other for rendering
 - In double buffering referred to as front and back buffers
 - Double buffering
 - To prevent flickering and other undesirable artifacts that will appear if an image that is currently being displayed is updated


Single Buffer



Double Buffering




Double Buffering




Buffers

- The depth-buffer (z-buffer)
 - Use for visibility determination
 - Tests can be enabled or disabled
 - Certain situations need to disable depth-buffer tests
 - e.g., when rendering translucent polygons



A simple three-dimensional scene



Z-buffer representation

Buffers

- The stencil buffer
 - General purpose buffer for doing things not possible with colour and depth buffer alone
 - E.g., for creating reflective surfaces and essential in some shadow rendering techniques, like shadow volumes
 - “Tag” pixels in one rendering pass to control their update in subsequent rendering passes
 - Can specify different rendering operations like
 - Stencil test fails
 - Stencil test passes & depth test fails
 - Stencil test passes & depth test passes

Buffers

- Clearing buffers

```
// set values to clear to  
glClearColor(r, g, b, a);  
glClearDepth(depth);  
glClearStencil(value);  
  
// clearing  
// for example, at the start of render_scene()  
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT |  
GL_STENCIL_BUFFER_BIT);
```


Buffers

- Multipass rendering
 - Render the scene multiple times
 - Composite the images in certain ways to achieve the desired effects in final scene
 - Two things to consider for each stage of a multipass algorithm
 - How to render the scene
 - How to composite the images


Buffers

- Stencil buffer for reflection

Without the stencil buffer




With the stencil buffer



Buffers

- Stencil buffer for reflection
 - Basic algorithm
 - Clear buffers
 - Disable depth buffer and draw mirror surface to stencil buffer
 - Enable depth buffer, draw reflected geometry where stencil buffer passes
 - Disable stencil buffer, draw normal scene



Compositing and Blending


- Alpha blending

- Alpha channel

- The 'A' component in the RGBA (or RGB α) colour mode


- α values range from 0.0 to 1.0

- Value of 1.0, surface completely opaque and blocks all incident light
 - Translucency of a surface given by $1.0 - \alpha$



Compositing and Blending

- Dealing with translucency in a physically correct manner is difficult due to
 - The complexity of the internal interactions of light and matter
 - Using a pipeline renderer



Compositing and Blending

- Blending equation
 - Can define source and destination blending factors for each RGBA component

$$\mathbf{s} = [s_r, s_g, s_b, s_a]$$

$$\mathbf{d} = [d_r, d_g, d_b, d_a]$$

Suppose that the source and destination colors are

$$\mathbf{a} = [a_r, a_g, a_b, a_a]$$

$$\mathbf{b} = [b_r, b_g, b_b, b_a]$$

Blend as


$$\mathbf{c}' = [a_r s_r + b_r d_r, a_g s_g + b_g d_g, a_b s_b + b_b d_b, a_a s_a + b_a d_a]$$

Compositing and Blending

- Hidden surface removal
 - Opaque polygons block all polygons behind them and affect the depth buffer
 - Translucent polygons should not affect depth buffer
 - Sort polygons first to remove order dependency
 - Easiest solution (not necessarily the best)
 - Render all opaque polygons first
 - Disable the depth test (or make the depth-buffer read-only)
 - Render translucent polygons in a back-to-front order



Mapping Techniques

- Billboarding
 - 2D texture map used as a 3D entity in the scene
 - Rotated so polygon's normal always faces the viewer
 - Creates illusion 2D image is 3D object




Mapping Techniques

- Billboard






Mapping Techniques

- Normal mapping
 - Textures used to perturb the surface normal
 - Actual geometry of surface does not change, just shaded as if it were different shape




Mapping Techniques

- Normal mapping
 - Outline looks smooth



+




=



Mapping Techniques

- Environment mapping
 - Quick but inaccurate way of generating effects like reflections (a.k.a reflection mapping)
 - Image of the surrounding environment used to project reflections onto object's surface
 - Common methods
 - Cube mapping
 - Spherical mapping

Mapping Techniques




➤ Cube mapping


- Popular because easily constructed
 - Place camera in centre of box and render the 6 different views
- Surrounding scene mapped onto surfaces of a cube

Mapping Techniques

- Environment mapping
 - Cube mapping





‘Unfolded’ cube



Mapping Techniques

- Environment mapping




Mapping Techniques

- Environment mapping



Visibility Determination



- Back-face detection
 - Fast and simple object-space approach for determining back face of a polygon
 - A polygon is facing forward if
$$-90^\circ \leq \theta \leq 90^\circ$$
 - or equivalently
$$\cos \theta \geq 0$$
 - Easier to test using dot product instead of cosine
$$\mathbf{n} \cdot \mathbf{v} \geq 0$$




Visibility Determination

- Back-face detection (cont.)
 - Right-hand system
 - Polygon is back face if $C \leq 0$
 - Default cull clockwise (CW) winding
 - Left-hand system
 - Polygon is back face if $C \geq 0$
 - Default cull counter-clockwise (CCW) winding

Polygon Winding




Right-handed System



Left-handed System

Visibility Determination

- Back-face detection (cont.)
 - Back-face culling
 - Removal of back-faced polygons
 - Usually happens before the clipping step in a rendering pipeline
 - Does this solve all visibility problems?



No.



Given these ‘front-facing’ polygons,
which sections are visible to the
viewer?

Visibility Determination

- The depth-sort method
 - Can be considered to be an object-space approach (the sorting part anyway)
 - General idea
 1. Sort polygons in order of decreasing depth (i.e. furthest to nearest)
 2. Scan convert polygons in order, starting from the furthest (greatest depth) polygon first
 - Nearer polygons will overwrite further polygons
 - This back-to-front rendering often referred to as the ‘painter’s algorithm’



Visibility Determination

- The depth-sort method (cont.)
 - Obviously, depth sorting has to occur before rasterization




Visibility Determination

- The depth-sort method (cont.)
 - This is easy if all the polygons' depths do not overlap
 - But, how do we sort polygons with overlapping depths?
 - For polygons with overlapping depth, need to make additional comparisons




Visibility Determination

- The depth-sort method (cont.)
 - Check whether the polygons overlap in x and y extents
 - Assuming normalized device coordinates where viewing direction is along z-axis, which is usually the case
 - If non-overlapping x and y, then the order does not matter




Visibility Determination

- The depth-sort method (cont.)
 - If previous test failed
 - Then test whether the all the vertices of one polygon lie on the same side of the plane defined by the other polygon
 - Order relative to viewing position




Visibility Determination

- The depth-sort method (cont.)
 - This still leaves two troublesome situations



Cyclic Overlapping




Piercing Polygons

Visibility Determination

- The depth-sort method (cont.)
 - In such cases, have to resort to splitting polygons
 - Does this solve all visibility problems?
 - Yes
- But depth sorting is somewhat laborious, particularly if we have to split polygons

Visibility Determination

- The depth-buffer method
 - Commonly used image-space approach
 - Compares surface depth values for each pixel position on the projection plane
 - Also referred to as the z-buffer method, since object depth is usually measured along z-axis
 - Easy to implement, in either hardware or software




A simple three-dimensional scene



Z-buffer representation

Visibility Determination

- The depth-buffer method (cont.)



Visibility Determination


- The depth-buffer method (cont.)
 - Other than colour buffer
 - Requires another buffer (as the name already implies) to store depth information
 - The depth buffer
 - Has the same resolution as the frame buffer and stores floating point numbers
 - Typically carried out in normalized coordinates, so depth values range from 0.0 (near clipping plane) to 1.0 (far clipping plane)

Visibility Determination

- The depth-buffer method (cont.)
 1. Initialise depth buffer and colour buffer for all positions
 - $\text{depthBuff}(x, y) = 1.0f$
 - $\text{colourBuff}(x, y) = \text{backgroundColour}$
 2. For each polygon in the scene (in any order)
 - For each projected (x, y) pixel position of polygon, calculate the depth z (if not already known)
 - If $z < \text{depthBuff}(x, y)$, then compute and set the surface colour at that position
 - $\text{depthBuff}(x, y) = z$
 - $\text{colourBuff}(x, y) = \text{surfColour}(x, y)$

Visibility Determination

- The depth-buffer method (cont.)



Visibility Determination

- The depth-buffer method (cont.)
 - Polygons can be processed in one at a time in any order and each polygon is processed one scan line at a time
 - Compatible with the pipeline architecture as it is performed polygon by polygon
 - Does this solve all visibility problems?

Yes

Used to be a problem when memory was expensive

References

- Among others, material sourced from
 - Hearn, Baker & Carithers, “Computer Graphics with OpenGL”, Prentice-Hall
 - Angel & Shreiner, “Interactive Computer Graphics: A Top-Down Approach with OpenGL”, Addison Wesley
 - Akenine-Moller, Haines & Hoffman, “Real Time Rendering”, A.K. Peters
 - Joey de Vries, “Learn OpenGL,” <https://learnopengl.com/>
 - Watt, “3D Computer Graphics”, Pearson/Addison Wesley
 - <http://en.wikipedia.org/wiki/>
 - IT OpenCourseWare, <https://ocw.mit.edu/>
 - http://http.developer.nvidia.com/GPUGems/gpugems_ch21.html